

Summary of Design Documents in nap-serv and pg-schemata

Overview

The two repositories contain a wide range of design documents. The `nap-serv` docs describe a multi-tenant project-costing system built on Node/Express and PostgreSQL. They cover architectural decisions, detailed module specifications (costing, accounting, inventory and SKU matching), coding conventions, API routing, authentication, testing practices and long-term roadmaps. The `pg-schemata` docs focus on best practices for the underlying Postgres schema layer and the release process for the associated NPM package.

nap-serv design docs

Architecture & tenancy

- **Tenant isolation.** Each customer (tenant) is stored in its own Postgres schema; queries always include a `tenant_id` column, and middleware injects the tenant context into requests ¹. The `Accounting` Module Design summary further emphasises strict intra-tenant isolation: all accounting data is scoped by `tenant_id` and there are no cross-tenant constraints, ensuring data remains isolated ².
- **Project/unit hierarchy.** Projects are subdivided into units, unit budgets, activities, cost lines and actual costs ³. This hierarchy supports fine-grained budgeting and cost tracking down to individual activities or change orders ⁴. All IDs use UUIDv7 and include a `tenant_id` for isolation ⁵.

Costing & accounting modules

- **Budgeting & cost tracking.** Budgets are versioned per unit and must be approved before any actuals can post ⁶. Cost lines represent planned amounts per activity and vendor; actual costs track real-world spend and cannot exceed the approved budget unless a change order or tolerance percentage allows it ⁷. Change orders capture approved adjustments to budgets and include workflow statuses (`pending`, `approved` or `rejected`) ⁸. The **cost-tracking** overview explains how projects, categories, activity budgets and cost lines connect, highlighting foreign-key relationships and sample SQL for variance reporting ⁹.
- **Actual cost logic.** Only project managers or automated imports can create actual cost entries. Default behaviors include auto-generated IDs, budget enforcement and pending approval status ¹⁰. An approval flow governs how pending costs become approved or rejected, with audit fields and tolerance rules ¹¹.
- **Accounting (General Ledger).** The accounting roadmap and task list describe an engine that automatically generates journal entries from cost lines, change orders, AP invoices, AR invoices and

inter-company transactions. It enforces balanced debits and credits and validates that postings occur within open fiscal periods ¹² .

- **Accounts Payable (AP) and Receivable (AR).** AP includes invoice approval workflows, cost-line integration, GL posting and vendor balance updates ¹³ . AR supports invoice creation, revenue recognition based on activity completion and posting of debit/credit entries to the ledger ¹⁴ .
- **Intercompany & consolidation.** The system supports companies within a tenant engaging in services for each other. Intercompany transactions auto-generate mirrored revenue and expense entries and include elimination tagging to facilitate consolidated reporting ¹⁵ . Later phases of the roadmap plan to generate consolidated P&L, balance sheets and elimination reports across companies ¹⁶ .

System features and roadmap

- **Feature overview.** The `FEATURES.md` document provides a design summary covering tenant isolation, activity codes, budgeting/actuals, billing models, GL, AP/AR, intercompany transactions, consolidation, chart of accounts, roles/permissions and reporting ¹ ¹⁷ . It distinguishes fixed-price, cost-plus and milestone billing and introduces **billing units**, abstract events to which costs and revenue are mapped ¹⁸ . The MVP includes tenant management, project costing, GL/AP, basic AR, intercompany logic, consolidated reporting and role-based security ¹⁹ . Post-MVP enhancements mention project scheduling, AI forecasting, change orders, retainage, multi-currency and external integrations ²⁰ .
- **Development roadmap.** `ROADMAP.md` defines seven phases: (1) core infrastructure and tenant isolation, (2) project costing engine, (3) GL/AP/AR scaffolding, (4) intercompany accounting, (5) consolidated reporting, (6) security & roles and (7) QA/hardening ²¹ ²² . A refactored `ROADMAPv2.md` notes that phases 1–4 are complete and lists remaining tasks for phase 5 (costing & accounting logic), phase 6 (consolidated reporting), phase 7 (security) and phase 8 (beta release) ²³ . It also includes an operational checklist and best-practice table showing that static module loading and routing have replaced deprecated dynamic loading ²⁴ .
- **Phase 5 tasks.** The phase-5 documents list specific tasks for activities, accounting, AP and AR modules. For activities, budget status must be enforced, remaining budget fields auto-calculated and change orders integrated ²⁵ . Accounting tasks include auto-posting journal entries, balanced transactions and linking entries to source documents ²⁶ . AP tasks involve invoice workflows, account mapping and vendor balance updates ¹³ , while AR tasks cover invoice creation, revenue recognition, postings and payment application ²⁷ . Shared logic requires a status engine, mapping tables and middleware to enforce fiscal rules ²⁸ .

API design & query parameters

- **Static routing.** `static-api-routing.md` shows how Express routers statically mount versioned module routes. The main `apiRoutes.js` mounts tenant and activities routers under `/api` ²⁹ , and module routers mount sub-routers under paths like `/v1/tenants` and `/v1/nap-users` ³⁰ . Tables map HTTP verbs to CRUD operations for each resource ³¹ .
- **Query parameter guidelines.** The frontend guide lists standard query parameters (`limit`, `offset`, `orderBy`, `columnWhitelist`, `includeDeactivated`) for list endpoints ³² . It explains cursor pagination using `cursor.<field>` ³³ , advanced condition-based filtering via `/where` ³⁴ , an endpoint for archived records ³⁵ and notes for developers such as using `encodeURIComponent` and not mixing offset and cursor pagination ³⁶ .

- **OpenAPI and reference docs.** The OpenAPI YAML defines query parameters accepted by the base controller: `limit`, `orderBy`, `includeDeactivated`, `columnWhitelist`, `JSON` `conditions` arrays and `cursor.{field}` values for cursor pagination ³⁷. The `/where` and `/archived` paths support `offset`, `limit`, `joinType` and field filters ³⁸. A companion reference summarises the allowed parameters and provides examples of simple, filtered and cursor-based requests ³⁹ ⁴⁰.

Coding conventions, testing and developer workflow

- **Contribution guidelines.** Contributors are asked to clone the repo, install dependencies and configure environment variables before running `npm start` ⁴¹. Bugs and feature requests should include clear descriptions and reproduction steps ⁴², while pull requests must follow a fork→branch→commit→PR workflow and include tests ⁴³. The document prescribes repository naming (plural for collections, singular for conceptual modules) and file naming conventions for repositories, controllers and routes ⁴⁴, enforces ESLint and Prettier styles ⁴⁵ and adopts Conventional Commits (`feat`, `fix`, etc.) ⁴⁶.
- **Coding conventions.** A separate document elaborates on the repository pattern, API setup, model definitions, controller patterns, routing and schema definitions. Repository objects export model classes using descriptive keys ⁴⁷, API routers aggregate sub-routers under versioned paths ⁴⁸ and controllers extend a base controller and are exported both as a default instance and a class for testing ⁴⁹. Table schemas are defined using `pg-schemata` typedefs with column and constraint definitions ⁵⁰.
- **Test setup.** Controller files should export both the class and its instance so tests can import the class directly ⁵¹. Constructors accept optional model injections for isolation ⁵², and tests should mock models before instantiating controllers ⁵³. The guide recommends avoiding `jest.mock` of the database, using factory helpers to create controllers, resetting mocks between tests, and writing custom unit tests for business logic ⁵⁴ ⁵⁵.

Authentication & authorization

- **Client and server requirements.** The frontend (React) needs the user's role for UI configuration and should automatically log out after 15 minutes of inactivity ⁵⁶. The backend authenticates via email and password and expects `email`, `user_name`, `tenant_code`, `role` and schema to be present in middleware ⁵⁷.
- **JWT structure.** A 15-minute access token and a 7-day refresh token are stored in secure, HTTP-only cookies ⁵⁸. An inactivity timeout is enforced by the frontend calling `/auth/refresh` periodically; if no request occurs, the user must reauthenticate ⁵⁹.
- **Backend middleware & endpoints.** The middleware verifies the token and attaches user and tenant context to the request ⁶⁰. `/auth/user` returns safe user profile details for UI rendering ⁶¹. A typical flow: user logs in, frontend fetches user info, periodically refreshes the session and logs out after inactivity ⁶². A summary table highlights token storage, expiry and safe exposure of user data ⁶³.

Additional nap-serv topics

- **Phase II activities summary.** The Phase II context explains the PERN stack, use of `pg-schemata` as ORM, UUIDv7 primary keys and multi-tenant architecture ⁵. It lists the modules—categories, activities, cost lines, activity budgets and actuals—and links to profitability views ⁶⁴. Project

management involves projects tied to clients and addresses ⁶⁵ . Next steps include scaffolding repositories, controllers and APIs for projects and budgets and writing tests ⁶⁶ .

- **Schema refactoring.** As the project evolved, a refactor plan outlines which schemas to keep (clients and cost lines), which to remove (activity budgets, actuals and project budgets) and which to refactor (projects, activities and profitability view) ⁶⁷ ⁶⁸ . New schemas introduced include units, unit budgets, actual costs and change orders ⁶⁹ .
- **SKU matching & onboarding.** A design document proposes automatic matching of vendor SKUs to catalog SKUs using vector embeddings. Models provide functions to find embeddings and similar records ⁷⁰ and to insert and query matches ⁷¹ . Matching logic compares vendor embeddings to catalog embeddings with a configurable threshold and auto-inserts high-confidence matches while returning low-confidence matches for user review ⁷² . A mermaid flowchart illustrates the end-to-end onboarding workflow: import vendor spreadsheets, validate data, generate embeddings, match to catalog SKUs, review matches and store confirmed matches; separate flows handle price updates, new/deprecated SKUs, maintenance and bid preparation ⁷³ .
- **Query model & OOP comparison.** An extensive guide compares object-oriented features in C++ and JavaScript, explaining class syntax, inheritance (`extends` vs. `:`), static methods, memory management, function overloading and private members ⁷⁴ ⁷⁵ . Summary notes highlight that JavaScript has class syntax and private fields but remains prototype-based and lacks true function overloading ⁷⁶ .

pg-schemata design docs

Design and best practices

- **Shared connection pool & tenant schemas.** The library uses a single `pg-promise` connection pool and switches schemas at runtime, avoiding the complexity of multiple pools ⁷⁷ . Each tenant has its own Postgres schema; models call `.setSchema()` after login to target the correct tenant ⁷⁸ .
- **Model schema definitions.** Models are defined as structured JavaScript objects with separate `columns` and `constraints` sections ⁷⁹ . UUIDs are used for primary keys and tenant IDs, generated with `gen_random_uuid()` ⁸⁰ .
- **ColumnSets and CRUD helpers.** The `TableModel` builds separate `insertColumnSet` and `updateColumnSet` to handle immutable and mutable fields, making inserts and updates safer ⁸¹ . It implements common read methods (`findById` , `findAll`) that automatically qualify tables with the current schema ⁸² .
- **Immutable fields & auto-creation.** The framework can mark columns immutable to prevent accidental updates ⁸³ . On tenant signup, schemas and tables are created programmatically ⁸⁴ .
- **Design principles.** Final principles emphasise a single connection pool, dynamic schema switching, structured model schemas, ubiquitous UUIDs, early construction of ColumnSets, dynamic read methods, careful handling of immutable fields and auto-generated SQL for migrations ⁸⁵ .

Release process

- **GitHub release.** To release a new version, the developer updates the `dev` branch, merges it into `main` with a release commit, bumps the version using `npm version` , then pushes `main` and tags to GitHub ⁸⁶ . A GitHub release can be drafted using the new tag and release notes ⁸⁷ .

- **npm publish.** After verifying login, run `npm publish` to publish the stable version, optionally using `--tag beta` for beta releases ⁸⁸. The document includes commands for managing npm dist-tags and notes on how to deploy documentation to GitHub Pages ⁸⁹.
- **Post-release sync.** After releasing, merge `main` back into `dev` to carry version bumps forward and push the updated `dev` branch ⁹⁰.

QueryModel conditions

- **Structured conditions.** The `QueryModel.findWhere()` method accepts an array of condition objects that map cleanly to SQL `WHERE` clauses. Basic equality is expressed as `{ column: value }` ⁹¹; `null` checks use `{ column: null }` to produce `IS NULL` ⁹².
- **Operators.** Supported operators include `$like` and `$ilike` for pattern matching ⁹³, `$from` and `$to` for range queries ⁹⁴, and `$in` for membership tests ⁹⁵. Nested `$or` and `$and` groups allow complex boolean logic ⁹⁶. Only these operator keys are supported; unsupported keys throw an error ⁹⁷.
- **Notes & future operators.** The document stresses that all values are parameterized to prevent SQL injection and hints at potential future operators like `$not`, `$gt`, `$lt`, `$between` and `$notIn` ⁹⁸.

Conclusion

The `nap-serv` design documents outline a robust, multi-tenant accounting platform with comprehensive budgeting, costing, accounting, inter-company transactions and reporting features. They provide detailed development roadmaps, coding conventions, testing practices, authentication strategies and specialized modules (e.g., SKU matching). The `pg-schemata` design docs complement this by prescribing best practices for Postgres schema management and model definitions, ensuring that the underlying data layer remains scalable, extensible and easy to maintain. Together, these documents offer a holistic view of building a SaaS-grade project accounting system using Node/Express and PostgreSQL.

¹ ¹⁵ ¹⁷ ¹⁸ ¹⁹ ²⁰ [raw.githubusercontent.com](https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/FEATURES.md)

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/FEATURES.md

² [raw.githubusercontent.com](https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/Accounting_Module_Design_Summary.md)

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/Accounting_Module_Design_Summary.md

³ ⁴ ⁶ [raw.githubusercontent.com](https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/budgeting_cost_tracking_spec.md)

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/budgeting_cost_tracking_spec.md

⁵ ⁶⁴ ⁶⁵ ⁶⁶ [raw.githubusercontent.com](https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/phase2_activities_summary.md)

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/phase2_activities_summary.md

⁷ ¹⁰ ¹¹ [raw.githubusercontent.com](https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/actual_costs_business_logic.md)

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/actual_costs_business_logic.md

⁸ [raw.githubusercontent.com](https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/change_order_lines_business_logic.md)

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/change_order_lines_business_logic.md

9 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/cost-tracking.md

12 13 14 25 26 27 28 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/Phase%205%20ROADMAP.md

16 21 22 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/ROADMAP.md

23 24 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/ROADMAPv2.md

29 30 31 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/static-api-routing.md

32 33 34 35 36 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/frontend-query-guide.md

37 38 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/query-parameters-openapi.yaml

39 40 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/query-parameters-reference.md

41 42 43 44 45 46 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/CONTRIBUTING.md

47 48 49 50 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/coding_conventions.md

51 52 53 54 55 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/test-setup.md

56 57 58 59 60 61 62 63 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/nap-auth-strategy.md

67 68 69 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/schema_refactor_plan.md

70 71 72 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/sku_matching_design.md

73 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/sku_onboarding_and_bidding_process.md

74 75 76 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/nap-serv/main/design_docs/cpp-vs-javascript-oop-comparison.md

77 78 79 80 81 82 83 84 85 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/pg-schemata/main/design_docs/DesignAndBestPractices.md

86 87 88 89 90 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/pg-schemata/main/design_docs/RELEASE.md

91 92 93 94 95 96 97 98 raw.githubusercontent.com

https://raw.githubusercontent.com/silverstone-i/pg-schemata/main/design_docs/querymodel-conditions.md